

Autoencoder reconstructions — known sounds rebuild cleanly, novel ones do not





known

known

novel

novel

input



rebuilt

